

Divide and Conquer DP

Table of Contents

- [Preconditions](#)
- [Generic implementation](#)
- [Things to look out for](#)
- [Practice Problems](#)
- [References](#)

Divide and Conquer is a dynamic programming optimization.

Preconditions

Some dynamic programming problems have a recurrence of this form:

$$dp(i, j) = \min_{k \leq j} \{ dp(i - 1, k) + C(k, j) \}$$

where $C(k, j)$ is some cost function.

Say $1 \leq i \leq n$ and $1 \leq j \leq m$, and evaluating C takes $O(1)$ time. Straightforward evaluation of the above recurrence is $O(nm^2)$. There are $n \times m$ states, and m transitions for each state.

Let $opt(i, j)$ be the value of k that minimizes the above expression. If $opt(i, j) \leq opt(i, j + 1)$ for all i, j , then

we can apply divide-and-conquer DP. This known as the *monotonicity condition*. The optimal "splitting point" for a fixed i increases as j increases.

This lets us solve for all states more efficiently. Say we compute $opt(i, j)$ for some fixed i and j . Then for any $j' < j$ we know that $opt(i, j') \leq opt(i, j)$. This means when computing $opt(i, j')$, we don't have to consider as many splitting points!

To minimize the runtime, we apply the idea behind divide and conquer. First, compute $opt(i, n/2)$. Then, compute $opt(i, n/4)$, knowing that it is less than or equal to $opt(i, n/2)$ and $opt(i, 3n/4)$ knowing that it is greater than or equal to $opt(i, n/2)$. By recursively keeping track of the lower and upper bounds on opt , we reach a $O(mn \log n)$ runtime. Each possible value of $opt(i, j)$ only appears in $\log n$ different nodes.

Note that it doesn't matter how "balanced" $opt(i, j)$ is. Across a fixed level, each value of k is used at most twice, and there are at most $\log n$ levels.

Generic implementation

Even though implementation varies based on problem, here's a fairly generic template. The function **compute** computes one row i of states **dp_cur**, given the previous row $i - 1$ of states **dp_before**. It has to be called with **compute(0, n-1, 0, n-1)**.

```

int n;
long long C(int i, int j);
vector<long long> dp_before(n), dp_cur(n);

// compute dp_cur[l], ... dp_cur[r] (inclusive)
void compute(int l, int r, int optl, int optr)
{
    if (l > r)
        return;
    int mid = (l + r) >> 1;
    pair<long long, int> best = {INF, -1};

    for (int k = optl; k <= min(mid, optr); k++)
        best = min(best, {dp_before[k - 1] + C

    dp_cur[mid] = best.first;
    int opt = best.second;

    compute(l, mid - 1, optl, opt);
    compute(mid + 1, r, opt, optr);
}

```

Things to look out for

The greatest difficulty with Divide and Conquer DP problems is proving the monotonicity of *opt*. Many Divide and Conquer DP problems can also be solved with the Convex Hull trick or vice-versa. It is useful to know and understand both!

Practice Problems

- [Codeforces - Ciel and Gondolas](#) (Be careful with I/O!)
- [SPOJ - LARMY](#)
- [Codechef - CHEFAOR](#)
- [Hackerrank - Guardians of the Lunatics](#)
- [ACM ICPC World Finals 2017 - Money](#)

References

- [Quora Answer by Michael Levin](#)
- [Video Tutorial by "Sothe" the Algorithm Wolf](#)

(c) 2014-2019 translation by <http://github.com/e-maxx-eng>

07:80/112